



National Technical
University of Ukraine
"Igor Sikorsky
Kyiv Polytechnic Institute"



Institute of
Physics and
Technology

Intellectual Data Analysis


Practice 5: Fully Connected Neural Nets

Dr. Nataliya K. Sakhnenko


Please review Lectures 6-7 before this practical

Multilayer Fully Connected (FC) ANNs

"Non-deep" feedforward neural network





Multilayer Perceptron (MLP)



Deep neural network

Biological Neuron versus Artificial Neural Network




<https://stats.stackexchange.com/questions/182734/what-is-the-difference-between-a-neural-network-and-a-deep-neural-network-and-w>

Why PyTorch

- Dynamic Computation Graph:** Unlike TensorFlow's static graphs, PyTorch builds the computation graph on the fly, making it more flexible and easier to debug.
- Pythonic:** PyTorch feels like native Python, making it more intuitive for Python developers.
- Strong Community Support:** PyTorch has a growing and active community with extensive documentation and resources.
- Seamless Integration:** It works well with popular Python libraries like NumPy and Pandas.

<https://docs.pytorch.org/tutorials/beginner/basics/intro.html>



1. [Tensors](#)
2. [Datasets and DataLoaders](#)
4. [Build Model](#)
5. [Automatic Differentiation](#)
6. [Optimization Loop](#)

Build Models with torch.nn

```
class NeuralNetwork(nn.Module):  
    def __init__(self):  
        super().__init__()  
        ...  
  
    def forward(self, x):  
        ...  
        return ...
```

```
model =  
    NeuralNetwork() .to(device)
```

✓ Loss Function

[nn.MSELoss](#) for regression tasks

[nn.CrossEntropyLoss](#) for classification

✓ Optimizer

optimizer = torch.optim.SGD([model.parameters\(\)](#), lr=learning_rate)

Optimization

Inside the training loop, optimization happens in three steps:

- Call `optimizer.zero_grad()` to reset the gradients of model parameters. Gradients by default add up; to prevent double-counting, we explicitly zero them at each iteration.
- Backpropagate the prediction loss with a call to `loss.backward()`.
- Once we have our gradients, we call `optimizer.step()` to adjust the parameters by the gradients collected in the backward pass.

```
def train_loop(dataloader, model, loss_fn,
               optimizer):

    model.train() # added for best practices
    for batch, (X, y) in enumerate(dataloader):
        optimizer.zero_grad()
        # Compute prediction and loss
        pred = model(X)
        loss = loss_fn(pred, y)
        # Backpropagation
        loss.backward()
        optimizer.step()
```

```
def test_loop(dataloader, model, loss_fn):

    model.eval()
    test_loss, correct = 0, 0
    with torch.no_grad():
        for X, y in dataloader:
            pred = model(X)
            test_loss += loss_fn(pred, y)
            correct += ...
```

A PyTorch Workflow

